

CHRIST (Deemed to be University)

Department of Computer Science and Engineering

Digital Ethics & Responsible AI Portfolio

Questions that make AI Answerable

CIA 3 – Component 2: Group Presentation & Digital Ethics Portfolio

Unit 5: Generative AI, Truth, Creativity, and Global Ethical Order

Course: Responsible AI

Section: 5BTCS B

Team Member	Register Number
Aman	2460324
Anushka Tater	2460336

Project Repository:

<https://github.com/aman27singh/digital-ethics-portfolio>

September 22, 2026

Contents

Abstract	2
1 Introduction	3
1.1 Project Overview	3
1.2 Scope: Unit 5 Coverage	3
1.3 Repository and Implementation	3
1.4 Team Composition	4
2 Methodology	5
3 Case Studies	6
3.1 Case 01 – Generative AI Digital Ethics Portfolio (Joint)	6
3.2 Case 02 – Khan Academy Data Audit & AI Evaluation	6
3.3 Case 03 – The Cartel That Never Spoke: Tacit Collusion Detection	6
3.4 Case 04 – Dutch Childcare Benefits Scandal (Toeslagenaffaire)	7
3.5 Case 05 – AI Chatbot in Banking (Joint)	7
3.6 Case 06 – Palantir’s AI in Military Decision-Making	8
3.7 Case 07 – Rx Sentinel: Fake Prescription Detection	8
3.8 Case 08 – Strava Responsible Data Audit	9
3.9 Case 09 – Kargu-2 Autonomous Drone Incident	9
4 Team Contributions	10
4.1 Aman (2460324)	10
4.2 Anushka Tater (2460336)	10
4.3 Joint Work	11
5 Findings and Conclusion	12
5.1 Cross-Cutting Findings	12
5.2 Conclusion	12
Project Repository	13

Abstract

This report documents the Digital Ethics & Responsible AI Portfolio developed by Aman (2460324) and Anushka Tater (2460336) for the CIA 3 Component 2 assessment of the Responsible AI course (Section 5BTCS B), Christ University. The portfolio addresses **Unit 5: Generative AI, Truth, Creativity, and Global Ethical Order** through nine case studies spanning synthetic media and misinformation, authorship and creative labor, algorithmic collusion in markets, AI in warfare and autonomous weapons, EdTech and child-data privacy, healthcare fraud detection, fitness-data surveillance, and public-sector algorithmic bias. Each case study applies recognized ethical frameworks – including utilitarian, deontological, rights-based, and care-ethics perspectives – together with practical audit tools such as the AI Ethics Canvas, to real-world deployments of generative and automated decision-making systems. The team’s overarching finding is that governance power over AI systems must be distributed and contestable: outputs should be labelled, verified, and independently audited, and meaningful human oversight must be designed into high-stakes systems before deployment rather than appended after harm occurs. The complete portfolio is implemented as a static web application (HTML/CSS/JavaScript) and is publicly available on GitHub.

1. Introduction

1.1 Project Overview

The *Digital Ethics & Responsible AI Portfolio* is a two-member group project submitted as part of CIA 3 Component 2 for the Responsible AI course. The project takes the form of an interactive web portfolio that curates nine case studies examining the ethical, societal, and governance implications of Generative AI and other intelligent, automated decision-making systems. The guiding principle of the work is captured in its own tagline: “*Questions that make AI answerable.*” Across every case study, the team places human dignity, accountability, and equity at the center of the analysis rather than treating ethics as an afterthought to technical capability.

1.2 Scope: Unit 5 Coverage

The portfolio was developed to address **Unit 5 – Generative AI, Truth, Creativity, and Global Ethical Order**. Table 1.1 summarizes how the six thematic areas of this unit map onto the nine case studies produced by the team.

1.3 Repository and Implementation

The portfolio is implemented as a static, client-side web application and hosted publicly on GitHub at <https://github.com/aman27singh/digital-ethics-portfolio>. The repository structure separates presentation from content:

- `index.html` – landing page containing the hero section, topic-coverage rubric, case-study gallery, and team section.
- `project.html` & `project-detail.js` – template and logic for individual case-study detail pages.
- `styles.css`, `overrides.css`, `project.css`, `originals.css` – styling layers.
- `script.js` – client-side interactivity (filtering, reveal animations).

Theme	Description	Ownership
Synthetic Reality & Truth	Evaluation of deepfakes, misinformation, and epistemic risk in AI-generated media and public discourse.	Joint / Anushka
Authorship & Creative Labor	Analysis of copyright, fair training-data compensation, and the moral status of machine-generated content.	Joint Portfolio
AI in Warfare & Propaganda	Ethical inquiry into autonomous lethal weapons (Kargu-2), target scoring (Palantir), and International Humanitarian Law compliance.	Aman
Tacit Collusion & Market Governance	Ethical AI model design for detecting algorithmic price coordination without explicit human agreement.	Anushka
AI Tutor Audit & Generative AI Evaluation	Khanmigo chat-visibility audit, child-privacy assessment, and prompt verification against generative AI hallucinations.	Anushka
Global Governance & Equity	Addressing Global North–South asymmetries, the digital divide, and the limits of universal ethical frameworks.	Joint / Aman

Table 1.1: Unit-5 topic coverage matrix

- `originals/` – source images, canvases, and handwritten artifact scans used across case studies.
- `original-viewer.html` – viewer for original hand-drawn/scanned artifacts.

The site allows visitors to filter the nine case studies by author (Joint, Aman, or Anushka) and links each case card to a detailed write-up.

1.4 Team Composition

Name	Register No.	Section
Aman	2460324	5BTCS B
Anushka Tater	2460336	5BTCS B

Table 1.2: Team members

Contact details listed on the portfolio:

- Aman: aman.bhagat@btech.christuniversity.in
- Anushka: anushka.tater@btech.christuniversity.in

2. Methodology

The team applied a mixed set of ethical-analysis tools consistently across the nine case studies:

- **Multi-framework ethical analysis:** utilitarian, deontological, rights-based, virtue, and care-ethics lenses were used, particularly for high-stakes military and financial AI systems.
- **AI Ethics Canvas:** a structured worksheet used to map stakeholders, risks, and mitigations for the Dutch Childcare Benefits Scandal and the Kargu-2 drone incident.
- **Data/privacy audits:** applied to consumer platforms (Khan Academy, Strava) to evaluate consent design, data collection scope, and surveillance risk.
- **Model design proposals:** original AI system designs (Tacit Collusion Index, Rx Sentinel) that embed human-in-the-loop safeguards rather than fully automated decisions.
- **Generative AI verification:** direct evaluation of AI tutor outputs (Khanmigo) for factual hallucinations, cross-checked against source material.

Work was divided between joint case studies (co-authored) and individually-led case studies, as reflected in the case-study index in Table 1.1 and the case-study chapter that follows.

3. Case Studies

3.1 Case 01 – Generative AI Digital Ethics Portfolio (Joint)

Authors: Aman & Anushka **Type:** CIA 3 Component 2, joint portfolio overview

A two-member group portfolio examining synthetic reality, deepfakes, authorship, AI in warfare, tacit collusion, and global AI governance.

- **Unit covered:** Unit-5: Generative AI, Truth, Creativity & Global Ethical Order.
- **Deliverables:** Ethical analysis, creative artifacts, guidelines, and a presentation plan.
- **Focus:** Synthetic reality, deepfakes, epistemic risk, and North–South equity.
- **Conclusion:** Governance power must be shared; AI outputs should be labelled, verified, and audited.

3.2 Case 02 – Khan Academy Data Audit & AI Evaluation

Author: Anushka Tater (2460336) **Type:** Data audit

An ethical audit of Khan Academy’s privacy practices, bundled consent, and a critical evaluation of generative AI prompt outputs.

- **Platforms examined:** Khan Academy, Khan Academy Kids, and the Khanmigo AI tutor.
- **Privacy risk:** Bundled consent at signup and teacher visibility into Khanmigo chat logs.
- **AI fact-check:** Five AI prompts were evaluated; hallucinations were identified regarding ad sales and dark patterns.
- **Finding:** Generative AI is a useful first-pass auditor but requires human verification before its claims are trusted.

3.3 Case 03 – The Cartel That Never Spoke: Tacit Collusion Detection

Author: Anushka Tater (2460336) **Type:** Model design

An ethical AI model design for detecting tacit algorithmic price coordination in housing and retail markets.

- **Challenge:** Pricing algorithms can coordinate on high prices without any direct human agreement.
- **Core metric:** A proposed *Tacit Collusion Index (TCI)* for real-time market supervision.
- **Ethical risk:** An accountability vacuum forms between software vendors and the landlords/retailers who deploy their pricing tools.
- **Mitigation:** Outcome-based algorithmic pricing audits and compulsory model retraining requirements.

3.4 Case 04 – Dutch Childcare Benefits Scandal (Toeslagenaffaire)

Author: Anushka Tater (2460336) **Type:** AI Ethics Canvas

An AI Ethics Canvas analysis of an automated fraud-scoring system that falsely accused approximately 26,000 families using discriminatory inputs.

- **Core issue:** Dual nationality was used as a fraud-risk factor, in violation of EU non-discrimination law.
- **Impact:** Roughly 26,000 families were financially ruined, and around 1,600 children were removed from their homes as a downstream consequence.
- **Failure mode:** No human review occurred before penalties were imposed; the scandal ultimately led to the resignation of the Dutch government cabinet.
- **Remedy proposed:** Ban demographic attributes as model inputs and mandate due-process human review before any punitive automated decision.

3.5 Case 05 – AI Chatbot in Banking (Joint)

Authors: Aman & Anushka **Type:** Card note / joint analysis

A joint ethical analysis evaluating moral duties and ethical standards across loan applications, deposits, and account-opening interactions handled by banking chatbots.

- **Scenarios covered:** Loan applications, deposit/credit assistance, and account opening.
- **Morality:** Applicants should be treated with dignity, offered unbiased options, and guided with care.

- **Ethics:** Systems must eliminate rich/poor bias, protect financial data privacy, and verify accounts rigorously.
- **Synthesis:** Empathetic interaction design must be combined with strict algorithmic fairness constraints.

3.6 Case 06 – Palantir’s AI in Military Decision-Making

Author: Aman (2460324) **Type:** Ethical analysis

An ethical analysis of AI-assisted military decision-making, target scoring, and the human-oversight gap in systems such as Palantir’s AIP.

- **Central question:** Can an opaque decision-support system recommend lethal action without eroding human accountability?
- **Frameworks applied:** Utilitarian, deontological, virtue, rights-based, and care ethics.
- **Focus areas:** Human oversight, system transparency, civilian protection, and compliance with International Humanitarian Law (IHL).
- **Conclusion:** Responsibility must be designed into a system before deployment, not appended after harm has already occurred.

3.7 Case 07 – Rx Sentinel: Fake Prescription Detection

Author: Aman (2460324) **Type:** Model design

An ethical AI model design for detecting fraudulent drug prescriptions while avoiding automated denial of legitimate care.

- **Problem:** Prescription fraud fuels substance addiction, but false positives risk blocking patients from needed medical care.
- **System output:** A risk score, a flag type, a document-anomaly map, and a human-review alert – never an automatic denial.
- **Safeguard:** Specialty-specific baselines are used for fields such as oncology and pain management, where legitimate prescribing patterns differ from the general population.
- **Governing rule:** The AI narrows the review queue; a qualified human clinician makes the final call.

3.8 Case 08 – Strava Responsible Data Audit

Author: Aman (2460324) **Type:** Data audit

A responsible-data audit of fitness tracking covering location privacy, surveillance risk, and digital inequality.

- **Data collected:** GPS routes, biometric signals, social content, and device-level signals.
- **Risk:** Aggregated activity data can reveal home addresses, daily routines, and other sensitive locations.
- **Gap identified:** Platform terms of service are bundled, and granular location-privacy settings require active discovery by the user.
- **Finding:** The Strava Metro product connects individual privacy risk, sampling bias, and municipal planning power in ways users are rarely made aware of.

3.9 Case 09 – Kargu-2 Autonomous Drone Incident

Author: Aman (2460324) **Type:** AI Ethics Canvas

An AI Ethics Canvas analysis of autonomy, meaningful human control, and the military responsibility gap, centered on the reported 2020 Kargu-2 loitering-munition incident in Libya.

- **Issue:** “Fire, forget, and find” loitering-munition behavior enables targeting decisions without a live data link to a human operator.
- **Stakeholders:** Manufacturers, military commanders, civilians, regulators, and the individuals targeted.
- **Risk:** Target misclassification, weapon proliferation, and erosion of the IHL principle of distinction between combatants and civilians.
- **Recommendation:** Require live authorization of each individual engagement and mandatory logging of black-box flight data for post-hoc accountability.

4. Team Contributions

4.1 Aman (2460324)

Specializes in military AI ethics, autonomous lethal weapons governance, healthcare fraud-triaging models, and fitness/location data-privacy audits.

Contributions:

- Palantir AI Case Study (Case 06)
- Rx Sentinel Fraud Design (Case 07)
- Strava Data Audit (Case 08)
- Kargu-2 Drone Canvas (Case 09)
- Banking Chatbot – Joint (Case 05)
- Global Governance & Equity theme – Joint contribution

4.2 Anushka Tater (2460336)

Focuses on EdTech data audits, AI tutor prompt verification, tacit algorithmic price-collusion detection models, and public-sector algorithmic bias analysis.

Contributions:

- Khan Academy Audit (Case 02)
- Generative AI Verification (part of Case 02)
- Tacit Collusion Design (Case 03)
- Dutch Childcare Canvas (Case 04)
- Banking Chatbot – Joint (Case 05)
- Synthetic Reality & Truth theme – Joint contribution

4.3 Joint Work

Both members co-authored the overarching Generative AI Digital Ethics Portfolio (Case 01), the Banking Chatbot analysis (Case 05), and the Global Governance & Equity theme addressing Global North–South asymmetries in AI regulation.

5. Findings and Conclusion

5.1 Cross-Cutting Findings

1. **Accountability gaps recur across domains.** Whether in military targeting (Palantir, Kargu-2), welfare fraud detection (Toeslagenaffaire), or algorithmic pricing (tacit collusion), the case studies repeatedly surface the same structural problem: automated systems make or heavily influence consequential decisions while no single party bears clear responsibility for the outcome.
2. **Human-in-the-loop design is a recurring safeguard.** The Rx Sentinel and banking-chatbot proposals both embed a rule that AI may narrow options or flag risk, but a human must make the final, high-stakes decision.
3. **Consent and transparency are frequently superficial.** The Khan Academy and Strava audits both found that consent is bundled and that privacy-protective settings are hard to discover, shifting the burden of protection onto users least equipped to exercise it.
4. **Generative AI needs verification layers.** The Khanmigo evaluation demonstrated that generative AI tools can hallucinate even in low-stakes educational contexts, reinforcing the need for human fact-checking before AI outputs are trusted or acted upon.
5. **Global equity is a governance issue, not just a technical one.** The Global Governance & Equity theme highlights that universal ethical frameworks can obscure Global North–South asymmetries in who sets AI rules and who bears AI’s costs.

5.2 Conclusion

Across all nine case studies, the team’s central and recurring conclusion is that AI systems – generative or otherwise – must be made *answerable*: their outputs should be labelled, their decisions verifiable, and their deployment subject to independent audit. Governance power over these systems should not be concentrated solely in the hands of the organizations that build or deploy them. Instead, meaningful human oversight, contestability, and accountability must be designed into high-stakes AI systems from the outset, rather than retrofitted after harm has already occurred.

Project Repository

The full interactive portfolio, including all visual artifacts, original scanned notes, and detailed per-case pages, is available at:

<https://github.com/aman27singh/digital-ethics-portfolio>