



# Palantir's AI in Military Decision-Making

*Ethical Case Study Analysis — Responsible AI*

---

**Student Name:** Aman

**Register No.:** 2460324

**Class:** 5BTCS B

# Background of the Case

02

## Organization

Palantir Technologies was founded in 2003 with initial funding from the CIA's In-Q-Tel. Named after the all-seeing stones in Tolkien's Lord of the Rings, it builds AI-powered data platforms for government and enterprise clients.

Key products include Palantir Gotham (used by intelligence agencies and military) and Palantir AIP (Artificial Intelligence Platform), deployed by the US Army, JSOC, NSA, FBI, and allied forces across 40+ countries.

## AI System Purpose

Palantir's military platforms aggregate massive volumes of intelligence data — surveillance feeds, signals intercepts, satellite imagery, and social media — to build real-time operational pictures of conflict zones.

The AI identifies persons of interest, maps insurgent networks, assigns threat risk scores, and surfaces targeting recommendations to military commanders, enabling faster decision-making in high-pressure battlefield environments.

## Timeline of Events

2003 — Founded; CIA In-Q-Tel funding  
2010s — Adopted by JSOC, NSA & CIA for counterterrorism  
2018 — Google exits Project Maven (ethical protest); Palantir steps in  
2019 — US Army awards Palantir major battlefield AI contract  
2022 — Deployed in Ukraine conflict to assist military decisions  
2023 — Palantir AIP launched; enables natural language battlefield queries  
2024 — Awarded \$480M US Army contract extension

*Palantir's systems are now operational in over 40 countries, including active conflict zones.*

## What Happened?

Palantir's AI systems, particularly Gotham and AIP, have been used by the US military and intelligence agencies to identify targets, predict insurgent behavior, and guide lethal drone strikes. In 2022, Palantir deployed its systems in Ukraine to assist battlefield decisions. Reports from human rights groups indicated civilian casualties linked to AI-assisted targeting decisions where human verification was minimal or bypassed under operational pressure.

## Key Ethical Issues

- Autonomous targeting without consent
- Opaque algorithmic decision-making
- Civilian harm & lack of accountability
- Erosion of meaningful human oversight
- Data sourced from mass surveillance

## Stakeholders Affected

### Civilians in Conflict Zones

Primary victims of AI-assisted targeting errors

### Military Operators

Reduced decision autonomy; moral offloading to AI

### Palantir Employees

Ethical concerns over product end-use (protests within company)

### Governments & NATO

Policy accountability for AI-enabled warfare

## Technology Used

- Machine learning & NLP for pattern detection
- Graph-based link analysis to map networks
- Predictive pattern recognition for threat scoring
- Real-time multi-source data fusion engine
- Drone & satellite feed integration (Project Maven)
- Natural language command interface (AIP)
- Computer vision for imagery analysis
- Anomaly detection across financial & comms data

## Data Sources

- NSA SIGINT (signals intelligence) intercepts
- Satellite & drone surveillance imagery
- Social media monitoring & sentiment analysis
- Interrogation reports & human intelligence (HUMINT)
- Financial transaction records & banking data
- Mobile phone metadata, call logs & GPS location
- Biometric databases (facial, voice, fingerprint)
- Open-source intelligence (OSINT) & news feeds

## Decision-Making Process

- Flags persons of interest via AI risk scores
- Builds social network maps of targets
- Ranks targets by predicted threat level
- Surfaces recommendations to operators
- Humans can approve or override (in theory)
- In practice: speed pressure limits oversight
- No explainability — black-box reasoning
- Lethal action possible within minutes of flag

*Palantir AIP operates via a chat-like interface allowing commanders to query battlefield data in natural language — raising questions about the speed and scrutiny of lethal decisions.*

## Does the AI maximize overall benefit?

### ✓ Who Benefits

- Faster threat identification reduces soldier casualties
- Processing vast intelligence faster than humanly possible
- Claimed reduction in collateral damage through precision
- Enables pre-emptive disruption of terror networks

### ✗ Who Is Harmed

- Civilian deaths from erroneous AI targeting (documented in Somalia, Syria)
- Algorithmic errors compounded by speed — no time to verify
- Long-term harm: normalizes autonomous lethal decisions
- Communities live under AI surveillance with no recourse

*Verdict: Utilitarian justification is weak — benefits accrue to militaries while harms fall disproportionately on civilians with no political voice.*

## IHL Violation Risk

International Humanitarian Law demands distinction between combatants and civilians. Palantir's probabilistic targeting cannot guarantee this distinction — violating a non-negotiable legal duty.

## Duty of Transparency

Operators and affected communities have a right to understand how decisions are made. Palantir's black-box model violates this duty — commanders cannot explain an AI's reasoning.

## Moral Agency Abdication

Deontology holds that moral actors must own their decisions. Delegating the choice to kill to an algorithm creates a 'responsibility vacuum' — no one is truly accountable.

## Fairness & Due Process

Individuals flagged by the AI as threats have no ability to contest, appeal, or even know they are being targeted — a fundamental violation of procedural justice.

# Ethical Analysis — Virtue Ethics

07

*Did Palantir and its leadership demonstrate the virtues expected of responsible technologists?*

| Virtue         | Expected Behavior                                                                                | Palantir's Actual Conduct                                                                                   | Verdict |
|----------------|--------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------|---------|
| Honesty        | Transparent disclosure of AI capabilities, limitations, and failure rates to client governments. | Palantir does not publicly disclose algorithmic accuracy or civilian harm incidents tied to its products.   | LACKING |
| Courage        | Refusing contracts that violate ethical norms — as Google did with Project Maven in 2018.        | Palantir actively sought Project Maven after Google's withdrawal, prioritizing revenue over ethical stance. | ABSENT  |
| Responsibility | Ensuring post-deployment oversight and mechanisms for correcting harm.                           | No public accountability framework exists for Palantir's role in combat decisions or civilian harm.         | ABSENT  |
| Prudence       | Careful, measured deployment — ensuring AI augments rather than replaces human judgment.         | AIP's chat interface is designed for speed, not deliberation — optimizing operational tempo over caution.   | PARTIAL |

## Right to Life

CRITICAL

AI-assisted strikes have caused civilian deaths. Every incorrect prediction that results in a strike is an unrecoverable violation of the most fundamental human right.

## Right to Privacy

SEVERE

Palantir aggregates phone metadata, financial data, and social media without consent from individuals who are then profiled as security threats.

## Right to Due Process

SEVERE

Targets are identified and acted upon without trial, legal review, or any opportunity to challenge the AI's classification — particularly acute for wrongly flagged individuals.

## Right to Autonomy

HIGH

Communities under AI surveillance have their movements, associations, and behaviors continuously monitored, fundamentally constraining their freedom without legal basis.

## Right to Remedy

HIGH

There is no mechanism for victims of AI-targeting errors — or their families — to seek accountability, compensation, or even acknowledgment.

*Care ethics asks: Are human relationships, well-being, and vulnerable populations placed at the center of this technology?*

## Vulnerable Groups Ignored

- Civilians in active conflict zones — Somalia, Yemen, Syria, Ukraine — are the primary unintended victims of AI targeting errors.
- Ethnic and religious minorities are historically over-indexed in intelligence databases, making AI profiling inherently biased.
- Children and non-combatants have no mechanism within Palantir's system to be flagged as protected persons.
- Grieving families of AI-strike victims have no access to information or justice.

## Human Relationships & Well-Being

- Operators report psychological distress — moral injury — from making life-or-death decisions mediated by an opaque AI system they do not fully understand.
- The system eliminates the face-to-face human judgment historically central to ethical warfare conduct.
- Communities under AI surveillance develop collective trauma and distrust of institutions.
- Care ethics demands empathy and contextual understanding — qualities fundamentally absent from Palantir's algorithmic design.

## Mandatory Human Authorization



No AI system should be permitted to recommend lethal action without explicit, documented human authorization at every step.

## Algorithmic Transparency



Palantir must publish audit reports on model accuracy, false positive rates, and civilian harm incidents — reviewed by independent bodies.

## International AI Weapons Treaty



Governments must negotiate binding international law on autonomous weapons systems — similar to the Chemical Weapons Convention.

## Civilian Protection Embedding



IHL principles (distinction, proportionality, precaution) must be hardcoded into targeting AI as non-negotiable constraints, not advisory flags.

*Key Lesson: When AI systems make irreversible decisions, the absence of accountability is not a gap — it is a design flaw. Responsibility must be built in, not bolted on.*