

# Ethical Case Study Analysis

## Palantir's AI in Military Decision-Making

Submitted by: Aman | Reg. No. 2460324 | 5BTCS B

Course: Responsible AI | CIA 1 – Component 1

Department of Computer Science & Engineering

Christ (Deemed to be University), Bengaluru

---

### 1. Introduction

On any given day, intelligence analysts in the United States military receive thousands of data points — satellite imagery, intercepted communications, financial records, and social media activity — far too much for any human team to process in time. Since the mid-2000s, Palantir Technologies has positioned itself as the answer to this problem, offering AI-powered platforms that aggregate, analyze, and surface actionable intelligence at machine speed. What began as a CIA-backed data analytics firm in 2003 has evolved into one of the most consequential — and contested — players in modern warfare.

Palantir's two primary platforms, Gotham and the newer Artificial Intelligence Platform (AIP), have been deployed by the US Army, the Joint Special Operations Command (JSOC), NATO allies, and most recently on the Ukrainian battlefield. These systems do not merely organize data — they recommend targets, predict insurgent behavior, and help commanders decide where to strike. The ethical implications of this are significant: when an algorithm's output can result in a missile being fired, the questions of fairness, accountability, and transparency stop being abstract and become matters of life and death.

This report critically examines Palantir's use of AI in military decision-making through five ethical frameworks — Utilitarian, Deontological, Virtue, Rights-Based, and Care Ethics — and evaluates the case against the six pillars of Responsible AI. The goal is not to argue whether Palantir's technology is effective, but whether it is ethical, and what needs to change if such systems are to be used responsibly.

### 2. Ethical Challenges

#### 2.1 Core Ethical Concerns

The central ethical problem with Palantir's military AI systems is the delegation of lethal decision-making authority — even partially — to an opaque algorithm. Unlike AI systems in healthcare or

finance, where errors are costly but often reversible, targeting errors in warfare are permanent. The following ethical concerns emerge from the case:

- **Opacity of decision-making:** Palantir's systems use machine learning models whose internal reasoning is not interpretable by the operators who act on their outputs. A commander approving a strike based on an AI-generated risk score cannot fully explain why that score was assigned.
- **Inadequate human oversight:** While Palantir maintains that humans retain final authority, investigative reports from the Intercept and Amnesty International have documented cases where the speed of AI-assisted operations effectively eliminated meaningful review. The system creates what researchers call a 'human-on-the-loop' rather than a 'human-in-the-loop' dynamic.
- **Surveillance-based data sourcing:** The AI draws on data collected through mass surveillance — phone metadata, location tracking, and social media monitoring — gathered without the consent of those being profiled, many of whom are civilians in regions already under political or military stress.
- **Civilian harm and absence of accountability:** When AI-assisted strikes result in civilian casualties, there is no established mechanism to determine whether the AI's recommendation was a contributing factor, making accountability structurally impossible.
- **Normalization of autonomous lethal systems:** Each deployment of Palantir's technology without significant incident normalizes the idea that AI can and should guide lethal force, setting a precedent that future, more autonomous systems will inherit.

## 2.2 Stakeholder Impact

**Civilians in conflict zones:** Populations in Somalia, Yemen, Syria, and Ukraine bear the highest risk from targeting errors. They have no voice in the deployment of these systems and no recourse when errors occur.

**Military operators:** Soldiers and intelligence personnel experience reduced moral agency — their professional judgment is increasingly subordinated to algorithmic outputs they did not design and cannot fully interpret. Research has documented increased rates of moral injury among drone operators, and AI-mediated targeting deepens this psychological distance from the consequences of decisions.

**Palantir employees:** Internal dissent at Palantir is real and documented. In 2020 and 2021, employee protests over the company's immigration enforcement contracts showed that a portion of the workforce is morally uncomfortable with certain applications of the technology they build. This creates ethical tension within the organization itself.

**Governments and policymakers:** Governments deploying Palantir's systems take on legal and political accountability for outcomes, even when the technical decision-making is effectively outsourced to a private company. This creates a novel accountability gap in international law.

**Palantir as a corporation:** By operating in conflict zones and building products that inform lethal decisions, Palantir has taken on an ethical responsibility that goes far beyond typical enterprise software companies — a responsibility that its current governance structures do not adequately reflect.

---

## 3. Analysis Using Ethical Theories

### 3.1 Utilitarian Ethics

Utilitarianism evaluates actions by their outcomes — specifically, whether they produce the greatest good for the greatest number. Applied to Palantir's military AI, the utilitarian case is contested on both sides.

Proponents argue that AI-assisted targeting is more precise than conventional intelligence methods, potentially reducing total casualties by identifying threats more accurately and quickly. The US military's public position has long held that precision-guided, intelligence-led operations reduce civilian harm compared to area bombardment. If Palantir's AI genuinely improves targeting precision, then on a strict cost-benefit calculation, it could be justified.

However, this argument has significant weaknesses. First, the empirical evidence for AI-driven harm reduction is not publicly available — Palantir's model accuracy rates are classified, and independent audits of civilian harm linked to AI-assisted strikes have not been permitted. What is documented — through investigative journalism and NGO reporting — is a pattern of civilian casualties in operations where AI-generated intelligence played a role. Second, the benefits of the technology accrue primarily to military operators and governments, while the risks fall disproportionately on civilian populations who have no political representation in the decision to deploy these systems. A utilitarian framework that ignores the distribution of benefits and harms, not just their aggregate, is

incomplete. On balance, the utilitarian justification for Palantir's current deployment is weak until its accuracy claims can be independently verified.

### **3.2 Deontological Ethics**

Deontological ethics, most associated with Kant's categorical imperative, evaluates actions by their adherence to universal moral duties rather than their consequences. Palantir's military AI raises multiple deontological violations.

Kant's categorical imperative asks whether a moral rule can be universalized without contradiction. The rule 'AI systems may make recommendations that result in lethal force without human operators fully understanding the basis of those recommendations' cannot be universalized — it would produce a world in which individuals are killed by processes that no one is meaningfully accountable for, which most moral traditions would regard as fundamentally unjust.

International Humanitarian Law (IHL) represents a codified form of deontological duty in armed conflict. IHL's core principles — distinction (between combatants and civilians), proportionality, and precaution — impose specific obligations on those who conduct military operations. Palantir's probabilistic risk-scoring does not and cannot guarantee the distinction principle: a person assigned a high risk score by an algorithm is not the same as a lawful military target identified through proper legal review. The speed and opacity of AI-assisted targeting creates conditions under which these legal duties can be de facto bypassed, even if they remain formally in place.

There is also a duty of transparency owed to operators, to affected communities, and to society. When commanders cannot explain the reasoning behind a targeting decision — because that reasoning lives inside a machine learning model — they cannot fulfill their duty to account for their actions. This is not merely a practical problem; it is a deontological one.

### **3.3 Virtue Ethics**

Virtue ethics asks not what rules were followed or what outcomes resulted, but whether the people and institutions involved demonstrated the character virtues expected of responsible moral actors — honesty, courage, justice, and practical wisdom.

Palantir's conduct in this regard is difficult to assess charitably. The company has consistently declined to publish accuracy rates, failure analyses, or civilian harm assessments for its military products. This is a failure of honesty — not individual dishonesty, but institutional opacity that makes accountability impossible.

The contrast with Google is instructive. When Google was approached to continue developing Project Maven — the Pentagon's AI-for-drone targeting initiative — employees protested and the company ultimately walked away from the contract, citing ethical concerns. Palantir stepped in to fill that gap, actively pursuing contracts that others had judged ethically problematic. From a virtue ethics perspective, this represents a failure of what Aristotle would call *phronesis* — practical wisdom — and of institutional courage.

This is not to suggest that Palantir's employees are individually without virtue. Many are, by all accounts, talented and well-intentioned engineers. But virtue ethics also applies to organizations and to the cultures they cultivate. An organization that builds lethal targeting systems without robust internal ethical review, that accepts contracts without public ethical deliberation, and that resists external scrutiny is not demonstrating the institutional virtues that the gravity of its work demands.

### **3.4 Rights-Based Ethics**

Rights-based ethics holds that individuals possess fundamental rights that cannot be overridden by utilitarian calculations or institutional convenience. Palantir's military AI implicates several categories of rights in ways that are difficult to reconcile.

The right to life is the most fundamental. Every person killed as a result of an AI-targeting error — every false positive — represents an irreversible violation of this right. Unlike most algorithmic errors, these cannot be corrected after the fact. The person misidentified as a threat cannot be compensated or restored.

The right to privacy is violated by the mass surveillance data collection that feeds these systems. Individuals whose phone records, financial transactions, and online behavior are harvested without consent have been deprived of informational autonomy — and that data is then used to make assessments about whether they constitute a security threat. This is a double violation: first the surveillance, then the consequential judgment based on it.

The right to due process — deeply embedded in both domestic legal systems and international human rights law — is structurally absent from AI-assisted targeting. There is no hearing, no legal review, no opportunity for the identified individual to contest the AI's classification. This is not a minor procedural concern; the entire legitimacy of state violence in liberal legal traditions rests on the idea that force is applied through lawful, accountable processes. AI targeting as currently practiced bypasses this legitimacy.

Finally, there is a right to remedy. Families of civilians killed in AI-assisted strikes have, in most documented cases, been unable to determine whether an algorithm contributed to the decision, let alone seek redress. This absence of remedy is not incidental — it is a structural consequence of deploying classified AI systems in environments with limited legal accountability.

### **3.5 Care Ethics**

Care ethics, developed by scholars including Carol Gilligan and Nel Noddings, grounds moral analysis in relationships, context, and the well-being of vulnerable people. It asks: who is caring for whom, and who is being left unprotected?

From a care ethics perspective, Palantir's technology fails at almost every dimension. The populations most affected by AI targeting — civilians in active conflict zones, ethnic and religious minorities historically over-indexed in intelligence databases, children who cannot be distinguished from adults in metadata analysis — are precisely the people who most need protection and who receive the least.

Care ethics also draws attention to the moral injury experienced by military operators. Soldiers and intelligence analysts who make decisions based on AI recommendations they do not fully understand, in operations whose outcomes they cannot fully trace, experience a form of moral disorientation that traditional frameworks of military ethics were not designed to address. The human relationships that make ethical warfare conduct possible — face-to-face judgment, contextual understanding, empathy — are systematically attenuated by AI mediation.

Perhaps most strikingly, Palantir's system is designed for operational tempo, not for care. Its core value proposition is speed — faster intelligence, faster targeting, faster decisions. But care ethics insists that the slowness of careful, contextual, empathetic judgment is not a weakness to be engineered away; it is a moral virtue. A system optimized for speed in life-or-death decisions is, from a care ethics standpoint, optimized for the wrong thing.

---

## **4. Responsible AI Evaluation**

### **4.1 Fairness**

Palantir's systems draw on datasets that reflect historical biases in intelligence collection — certain ethnic communities, geographic regions, and political affiliations are disproportionately represented in surveillance databases. Machine learning models trained on this data will replicate and potentially

amplify these biases. There has been no independent audit of Palantir's military AI for demographic or geographic bias. The fairness criterion is not met.

## **4.2 Accountability**

There is no clear accountability structure for AI-assisted targeting errors. The private company that built the system, the government that deployed it, and the individual operators who acted on its outputs each have grounds to point elsewhere when harm occurs. This 'responsibility gap' — identified by AI ethicists including Luciano Floridi — is one of the most serious structural failures in the current deployment model. Accountability is not met.

## **4.3 Transparency**

Palantir's model architectures, training data, accuracy rates, and failure modes are classified or proprietary. Operators cannot explain the basis of AI-generated risk scores; affected communities have no visibility into the systems that profile them; independent researchers cannot audit the technology. Transparency is fundamentally absent.

## **4.4 Privacy**

The data sourcing practices underlying Palantir's military AI — mass surveillance, bulk metadata collection, cross-database aggregation — are fundamentally incompatible with the privacy rights of the individuals being profiled. Even if each individual data source were lawfully obtained, the combination of sources creates privacy violations that no single element would produce. Privacy is not adequately protected.

## **4.5 Safety**

The system's safety cannot be meaningfully assessed without access to accuracy data and error rates. What is known from open-source reporting is that AI-assisted operations have been associated with civilian casualties in multiple theaters. The absence of a mechanism for detecting and correcting AI errors in real time, combined with the irreversibility of lethal force, means that safety requirements are structurally unmet.

## **4.6 Human Oversight**

While Palantir and the military assert that humans retain final authority over targeting decisions, operational realities undermine this claim. When an AI system produces a recommendation under time pressure, and operators have limited time and limited interpretive access to the AI's reasoning, the 'human in the loop' becomes a rubber stamp rather than a genuine check. Meaningful human oversight requires the ability to understand, question, and override AI recommendations — capabilities that the current system design does not fully support.

---

## **5. Recommendations**

The ethical failures identified in this analysis are not inevitable features of military AI; they are correctable design and governance failures. The following recommendations address the core responsibility gaps:

### **5.1 Mandatory Human Authorization Protocol**

No AI-generated targeting recommendation should result in lethal action without explicit, documented human authorization at multiple levels of command. This is not merely a policy guideline — it should be a hard technical constraint embedded in the system, with an auditable record of every decision and the human who authorized it. This addresses the accountability gap and strengthens the human oversight criterion.

### **5.2 Mandatory Independent Algorithmic Auditing**

Before deployment and at regular intervals during deployment, Palantir's military AI systems should be subject to independent audits — conducted by parties with no commercial relationship to Palantir — assessing accuracy rates, demographic bias, false positive and negative rates, and documented civilian harm. Results should be shared with relevant government oversight bodies and, where possible, with the public. This addresses transparency and fairness failures.

### **5.3 International Legal Framework for Autonomous Weapons**

The international community must develop binding legal standards for AI use in armed conflict, analogous to the Chemical Weapons Convention. Such a framework should set minimum standards for human oversight, prohibit fully autonomous lethal systems, establish state responsibility for AI-related civilian harm, and create mechanisms for civilian redress. This addresses the rights-based violations and the absence of meaningful accountability.



## 5.4 IHL Compliance by Design

The principles of International Humanitarian Law — distinction, proportionality, and precaution — must be embedded as non-negotiable technical constraints in any AI system used in armed conflict, not as advisory outputs that operators can override. Palantir should be required, as a condition of government contracting, to demonstrate how these principles are implemented at the model level.

## 5.5 Whistleblower Protections and Internal Ethics Structures

Palantir employees who raise ethical concerns about specific contracts or deployments must have protected channels to do so without professional retaliation. Independent ethics boards — with genuine authority to review and potentially block certain contracts — should be established within the company, with representation from ethicists, human rights professionals, and affected communities where feasible.

---

## 6. Conclusion

Palantir's AI in military decision-making represents a genuinely novel ethical challenge — one that does not fit neatly into existing frameworks for either corporate responsibility or military ethics. The company occupies an unprecedented position: a private entity whose commercial products influence decisions about who lives and who dies in armed conflict, operating under minimal public scrutiny and without a coherent accountability structure.

The analysis across all five ethical theories points to a consistent conclusion: the current deployment model is ethically insufficient. It fails utilitarian standards because its benefits are unverified and its harms fall on the most vulnerable. It violates deontological duties of transparency, due process, and IHL compliance. It reflects institutional virtues of expediency over responsibility. It infringes on fundamental rights to life, privacy, due process, and remedy. And it embodies precisely the opposite of care ethics — a system optimized for speed rather than for the careful, contextual human judgment that life-or-death decisions demand.

None of this is an argument that AI has no place in military operations. Used appropriately — to process intelligence, identify patterns, reduce cognitive load — AI can genuinely support better decision-making. But 'better decision-making' in warfare means more ethical, more accountable, and more protective of human life, not faster or more efficient at lethal action.

The key lesson from this case is deceptively simple but genuinely important: when AI systems make irreversible decisions, the absence of accountability is not a gap to be patched later. It is a design flaw that must be addressed before, not after, deployment. Responsibility must be built into these systems from the ground up — legally, technically, and ethically. Until that standard is met, the ethical use of AI in military decision-making remains an aspiration rather than a reality.

---

*— End of Report —*