

RESPONSIBLE AI CIA-1 COMPONENT-2 RESPONSIBLE DATA AUDIT REPORT

Platform Audited: Khan Academy (khanacademy.org)

Course	Computer Science & Engineering — Responsible AI
Submitted By	Anushka Tater
Register No.	2460336
Department	Computer Science and Engineering
Class	5 BTCS B
Date of Submission	04 July 2026
AI Tool(s) Used	ChatGPT

Section 1: Introduction

1.1 Name of Selected Application

For this audit, I chose Khan Academy (khanacademy.org) — the free online learning platform — along with a look at its two related products, Khan Academy Kids and the Khanmigo AI tutor, since both feed into the same overall privacy ecosystem.

1.2 Reason for Selection

I picked Khan Academy for a few reasons, some practical and some personal:

- It's a platform almost everyone in engineering has used at some point — for brushing up on math, or just exploring a topic outside the syllabus — so its data practices actually matter to people like us, not just to some abstract 'user.'
- It's used heavily by school-going children, which immediately raises the stakes on privacy and consent — kids can't really read a privacy policy and decide for themselves, so how that consent is handled becomes a genuinely interesting ethical question.
- It recently added Khanmigo, an AI tutor built on large language models. That's a fairly new category of data (your actual conversation with an AI) and I wanted to see how a platform is handling that compared to the more familiar 'we collect your email' type of clause.
- Khan Academy is a nonprofit and explicitly says it doesn't sell data or run ads — which made me curious whether that claim actually holds up once you dig into the fine print, or whether it's mostly marketing language.

1.3 Brief Description of the Platform

Khan Academy is a 501(c)(3) nonprofit that offers free video lessons, practice exercises, and progress tracking across subjects like math, science, programming, economics, and the humanities. Individual students use it on their own, but a large chunk of its usage actually comes through schools — teachers set up classroom accounts and monitor student progress through it. Alongside the main site, Khan Academy Kids is aimed at younger learners and is managed almost entirely by parents or teachers, while Khanmigo is the newer AI tutor/teaching-assistant feature, which brings in a whole new layer of data (chat logs) and a different set of safety questions than a normal ed-tech app.

Section 2: Privacy and Consent Analysis

2.1 Data Collected

Going through the actual Privacy Policy and the Security Practices page (not just relying on the AI summary), here's what Khan Academy collects:

- Contact and profile information — name, email address, username, password, and profile details provided at sign-up.

-
- Account and authentication information — login credentials, sign-in method (including via Google/Facebook/Clever/school SSO).
 - Location information — general/geographic information inferred from device or network.
 - Browser and device information — device type, operating system, browser type, and similar technical identifiers.
 - Usage and learning-activity data — videos watched, exercises attempted, time spent, progress and mastery data, and (for Khanmigo) conversational/chat logs with the AI tutor.
 - Information collected from others — data provided by a parent, teacher, or school on behalf of a student, and information from connected third-party applications.
 - Donor and applicant information — for individuals who donate to or apply to work with the nonprofit.
 - Cookies and similar technologies — strictly necessary cookies (required) plus optional analytics and marketing cookies.

2.2 Purpose of Collection

The policy explains the 'why' behind most of this fairly clearly:

- Running and personalizing the actual learning experience — e.g., suggesting what to practice next based on your progress.
- Keeping users in the loop about their account, product changes, and (if you don't opt out) promotional emails.
- Letting users take part in special programs Khan Academy runs jointly with partner organizations.
- Internal research and analytics to improve the product — mostly framed as using de-identified or aggregated data rather than data tied to one person.
- For schools specifically, generating progress reports for teachers, under Data Protection Agreements that are supposed to align with FERPA.

2.3 Consent Mechanisms

These are the below given consent mechanisms observed:

- Signing up means agreeing to the Terms of Service and Privacy Policy in one go — a single click-wrap checkbox, not separate consent for each individual use of your data.
- Cookies are handled a little better — 'strictly necessary' cookies are non-negotiable, but analytics and marketing cookies can actually be turned off through a preferences panel.
- For kids under 13, the child obviously isn't the one consenting — it's either a parent giving verifiable consent, or, more commonly in school settings, the school itself consenting on the child's behalf under education-privacy law.
- If Khan Academy plans a material change to how it handles Student Personal Data, it's supposed to notify the School first and give them a choice before rolling it out — a decent safeguard.

2.4 User Rights and Controls

- You can generally view, edit, or delete your profile information from account settings.
- Parents get the ability to change a child's username within Khan Academy Kids.
- Depending on where you live (GDPR in the EU, CCPA/CPRA in parts of the US, or India's DPDP Act), users can typically request access to, correction of, or deletion of their data — though how smooth that process actually is wasn't something I could verify just from the policy text.
- Khan Academy repeats, in more than one place, that it doesn't sell personal data, and it's signed the Student Privacy Pledge, which is meant to stop targeted advertising to students specifically.

Screenshot: Khan Academy Privacy Policy page



Table of Contents

[Khan Academy Privacy Principles](#)
[Overview](#)
[Contact Khan Academy](#)
[Collection of information](#)
[Use of information](#)
[Disclosure of information](#)
[Retention of Personal Information](#)
[Sponsorship and advertising](#)
[Transparency and your choices](#)
[Your account settings](#)
[Security of your personal information](#)
[Links to third parties](#)
[Changes to this Privacy Policy](#)
[Additional privacy notices](#)
[Schools and student use](#)
[Children's Privacy Notice](#)
[Donor information](#)
[Paid products](#)
[Applicant information](#)
[Regional information](#)

Khan Academy Privacy Principles

Our Privacy Policy is fueled by our commitment to these Privacy Principles:

1. We're deeply committed to creating a safe and secure online environment for you.
2. We do not sell your personal information to third parties. We established Khan Academy as a nonprofit organization so that our mission of education and your trust will not be in conflict with a for-profit motive.
3. We strive to provide you with access to and control over the information you give us, and we take the protection of your information very seriously.
4. We take extra precautions for our younger learners as described in our [Children's Privacy Notice](#), including restricting child accounts to automatically block features that would allow a child to post or disclose personal information.
5. We do not display third party advertising on Khan Academy. Our mission is to provide you with a world class education, not to sell you products.

Overview

Khan Academy, Inc. ("Khan Academy", "our", "us", or "we") is a US-based charitable nonprofit organization. As a 501(c)(3) nonprofit, we use the information we collect to provide you with a better experience and fulfill our mission of providing a free, world-class education for anyone, anywhere. We established ourselves as a nonprofit organization so that a for-profit motive will not interfere with our mission of providing a trusted educational resource.

We are committed to protecting your privacy. We understand how important privacy is to you, and we are committed to protecting your privacy and to creating a safer environment for learners of all ages. This Privacy Policy explains how information is collected, used, disclosed, and protected by Khan Academy when you use (or access) our website at www.khanacademy.org, our mobile application, and

Section 3: Ethical Data Audit

3.1 Privacy Concerns

- Because consent is all bundled together at signup, it's genuinely hard for a student to know which specific piece of data is being used for which purpose .
- Khanmigo chats being visible to teachers and parents is probably the most interesting tension I found. It makes sense as a safety net, but it also means a student can't really ask the AI a sensitive or embarrassing question without an adult potentially seeing it later.
- Device, browser, and location data get collected in the background, and while that's fairly standard across the industry, the app itself doesn't really explain this to a regular user in plain language — one has to go digging in the policy.

3.2 Manipulative Practices or Dark Patterns

- No aggressive dark patterns were observed see on shopping apps — no countdown timers, no buried 'delete account' buttons. That tracks with it being a free, nonprofit platform with nothing to upsell.
- That said, the single accept-everything checkbox at signup is still a mild version of a 'take it or leave it' design — you can't say yes to the app but no to a specific data use.
- The streaks, badges, and notification nudges are clearly there to keep students coming back. That's understandable from a learning-motivation angle, but it's still a persuasive-design technique, and it's worth asking how that lands on younger, more impressionable users.

3.3 Data Power and Surveillance Concerns

- In school setups, teachers and admins can see a lot about what an individual student is doing — including AI chat logs — which creates a real power imbalance between the institution and the student being watched.
- Because Khan Academy sits on so much aggregated learning data, it's quietly in a position to influence what 'good progress' even looks like across huge numbers of students — that's a kind of soft power that doesn't get talked about much.
- The platform runs on Google Cloud, so even with Khan Academy's own encryption layered on top, the data ultimately passes through infrastructure that isn't fully in Khan Academy's hands.

3.4 Digital Inequality Issues

- Free doesn't mean equally accessible. One still needs a working device and a stable internet connection to actually use it, so students in low-connectivity areas — which is a real issue in parts of India — are left out regardless of cost.
- Since the personalization gets better the more you use the platform, students who can log in consistently probably get a noticeably better experience than students who can only access it occasionally — the gap doesn't close, it just moves.

-
- The whole school-account model assumes the school already has decent digital infrastructure in place. A well-funded school can negotiate proper data protection terms; an under-resourced one may not .

3.5 Representational Harm Concerns

- Khanmigo's responses come from an underlying language model, and those models are known to carry biases from their training data — so the examples, names, or scenarios it generates could unintentionally skew toward certain cultures or leave others out.
- If the 'educational research' analytics aren't anonymized carefully, there's a risk of accidentally building stereotyped profiles of certain schools or regions based on their performance numbers.
- I couldn't find any actual reported incident of this happening at Khan Academy — this is more of a structural risk that comes with using AI-generated content at scale, not a confirmed problem.

Section 4: Critical Evaluation of Generative AI

Five prompts were run through a Generative AI tool and then went back and fact-checked each answer against the real privacy policy . The full row-by-row comparison is in the Verification Table below; here's the summary of what I noticed.

4.1 Correct Observations Generated by AI

- It got the basic data categories right — name, email, usage/progress data, cookies — nothing surprising here, this is the kind of thing AI tools are usually reliable on.
- It correctly picked up on the Khanmigo chat-visibility issue, which I thought was a genuinely good catch and matched what the Help Center actually says.
- It correctly flagged that children's data needs extra protection and parental/school consent, which lines up with the actual policy.
- It gave Khan Academy credit for its security practices (encryption, SOC 2 audits) — and this checked out against the Security Practices page.

4.2 Missing Ethical Concerns

- It didn't bring up the digital-divide angle on its own — I had to specifically ask about digital inequality before it engaged with the idea that 'free' doesn't mean 'accessible to everyone.'
- It barely touched on the power imbalance between teachers/admins and students in terms of who can see whose data.
- It didn't mention algorithmic bias in Khanmigo's AI-generated responses at all until I asked directly — this felt like a real gap given how central the AI tutor now is to the platform.

4.3 Hallucinations or Incorrect Claims

- It confidently claimed Khan Academy sells data to advertisers and shows targeted ads — this is just wrong, and it directly contradicts what Khan Academy states in its own policy and its Student Privacy Pledge commitment.

-
- It also claimed there were commercial-style dark patterns like fake urgency or a hidden unsubscribe button — I couldn't find any evidence of this anywhere on the platform. It looked like the AI was just pattern-matching to what it usually sees in ad-funded apps rather than actually checking Khan Academy specifically.

4.4 Biases and Limitations of the AI Tool

- The AI seems to default to generic ed-tech/commercial-app assumptions instead of grounding its answers in what Khan Academy's policy actually says — it's pattern-matching, not fact-checking.
- Its knowledge has a cutoff, so newer changes (like anything Khanmigo-specific added in 2025–2026) might not be reflected accurately, which is exactly why I had to double check everything against the live policy page.
- It states things confidently whether they're right or wrong — there's no difference in tone between a verified fact and a guess, which is honestly the most dangerous part if you're not double-checking.
- It's decent at summarizing what a policy literally says, but noticeably weaker at reasoning about deeper stuff like power dynamics, inequality, or bias — those needed me to think it through myself rather than rely on the AI.

Prompt 1:

I am auditing Khan Academy (khanacademy.org) for a Responsible Data Ethics assignment. Please summarise what personal data Khan Academy collects from its users, why it collects that data, and whether it shares or sells that data to third parties. Base your answer on what you know about the platform's privacy policy.

Response:

According to Khan Academy's privacy policy, the platform collects several categories of personal data to provide and improve its educational services. This includes contact and profile information such as name, email address, username, and password; account authentication details including sign-in methods; browser and device information; general location information inferred from the user's device or network; cookies and similar technologies; learning and usage data such as videos watched, exercises completed, progress, mastery levels, and time spent on the platform. For Khanmigo, it also collects AI conversation logs. In addition, information may be provided by parents, teachers, schools, connected third-party applications, donors, or job applicants.

Khan Academy collects this information to create and manage user accounts, personalize learning experiences, recommend appropriate content, monitor student progress, communicate important account and product updates, support educational programs with partner organizations, conduct internal research and analytics (primarily using aggregated or de-identified data), and provide teachers with student progress reports under applicable education privacy laws.

The privacy policy states that Khan Academy **does not sell users' personal information** and **does not display third-party targeted advertisements**. While it shares data with certain service providers and educational partners when necessary to operate its services, these disclosures are for service delivery, security, legal compliance, or school administration—not for selling personal data or advertising purposes. This commitment is reinforced by Khan Academy's participation in the **Student Privacy Pledge**.

 Khan_Academy_Responsible_Data_A...



Prompt 2:

For the same Khan Academy audit: analyse the consent mechanisms the platform uses. How does a regular user consent to data collection? Are there separate opt-ins for different types of data, or is consent bundled? How does the platform handle consent for users who are children under 13? Identify any legal frameworks (e.g. COPPA, FERPA, GDPR) that apply.

Khan Academy primarily relies on **bundled consent** for personal data collection. During account creation, users agree to both the **Terms of Service** and the **Privacy Policy** through a single click-wrap agreement. This means users accept the platform's overall data collection and processing practices at once, rather than providing separate consent for individual categories of personal data.

However, consent is **not entirely bundled**. The platform distinguishes between different types of cookies. **Strictly necessary cookies** are required for the platform to function and cannot be disabled, while **analytics and marketing cookies** can be managed through a cookie preferences panel, giving users some control over optional tracking technologies.

For **children under 13**, Khan Academy follows child privacy regulations by requiring consent from a **parent or legal guardian**. In educational settings, schools may also provide consent on behalf of students under applicable education privacy laws. This school-mediated consent model allows teachers or school administrators to create and manage student accounts while complying with legal obligations. If Khan Academy plans to make significant changes to how Student Personal Data is handled, it states that schools will be notified and given the opportunity to review those changes before implementation.

Several legal frameworks influence Khan Academy's consent practices:

- **COPPA (Children's Online Privacy Protection Act)**: Requires verifiable parental consent before collecting personal information from children under 13.
- **FERPA (Family Educational Rights and Privacy Act)**: Governs the handling of student education records and allows schools to manage student data under educational agreements.
- **GDPR (General Data Protection Regulation)**: Provides users in the European Union with rights such as access, correction, deletion, and greater control over personal data.
- **CCPA/CPRA (California Consumer Privacy Act/California Privacy Rights Act)**: Grants California residents rights over their personal information, including access and deletion.
- **India's Digital Personal Data Protection (DPDP) Act**: Gives users rights to access, correct, and request deletion of their personal data where applicable.

Prompt-3

does the platform employ any manipulative design practices or 'dark patterns'? Look for things like fake urgency, hidden unsubscribe or account-deletion options, pre-ticked consent boxes, or deceptive framing. Also consider softer persuasive techniques such as streaks, badges, or notification nudges — do these raise any ethical concerns given that many users are school-age children?

Khan Academy does not appear to use aggressive manipulative design practices or major dark patterns commonly found in commercial platforms. There is no evidence of tactics such as fake urgency, countdown timers, artificial scarcity, misleading purchase pressure, hidden unsubscribe buttons, or intentionally difficult account-deletion processes. This aligns with Khan Academy's nonprofit model, where the platform does not rely on advertisements or maximizing user purchases.

However, some areas raise ethical questions from a responsible design perspective. During account creation, consent is mostly handled through a single Terms of Service and Privacy Policy acceptance, meaning users agree to the overall data practices together instead of choosing separate permissions for each type of data use. While this is a common industry practice, it creates a mild “take it or leave it” consent structure because users cannot fully customize what data processing they accept.

Khan Academy also uses softer persuasive design techniques such as learning streaks, badges, progress tracking, achievements, reminders, and notification nudges. These features are designed to encourage motivation, consistency, and positive learning habits rather than increase commercial engagement. However, because many Khan Academy users are children and students, these techniques still deserve ethical attention.

For younger users, gamification may create pressure to maintain streaks or complete activities mainly for rewards rather than genuine learning. Continuous progress tracking may also influence how students perceive success, especially when teachers or parents monitor performance data. Therefore, even positive educational nudges should be designed carefully to avoid creating stress, unhealthy comparison, or excessive dependence on external rewards.

Overall, Khan Academy does not show evidence of harmful dark patterns, but its bundled consent approach and gamified engagement features represent softer forms of persuasive design that should be continuously evaluated because of the platform's large school-age user base.

sPrompt -4:

: evaluate whether the platform raises any digital-inequality or equity concerns. Does 'free access' truly mean equal access for all learners globally? Also consider representational harm — could the platform's AI features (such as the Khanmigo AI tutor built on a large language model) introduce bias, cultural blind spots, or unequal learning experiences for students from different backgrounds?

Although Khan Academy provides **free access** to educational content, free availability does not automatically guarantee **equal access for all learners globally**. The platform removes financial barriers by not charging students for lessons, but other forms of inequality still affect who can benefit from it.

A major concern is the **digital divide**. Students still require a suitable device, reliable internet connectivity, electricity, and digital literacy skills to consistently use Khan Academy. Learners in rural areas, low-income households, or regions with poor internet infrastructure may not receive the same benefits as students with constant access to high-speed connections and personal devices. This means that even though the platform is free, unequal technological access can create unequal learning opportunities.

Another equity concern comes from **personalized learning systems**. Khan Academy improves recommendations based on user activity, progress history, and learning data. Students who can use the platform regularly generate more data, allowing the system to personalize their experience more effectively. In contrast, learners with limited access may receive a less optimized experience because the system has less information about their learning patterns.

Khanmigo, Khan Academy's AI tutor, introduces additional concerns around **representational harm and algorithmic bias**. Since Khanmigo is built using a large language model, its responses may reflect biases present in the data used to train such AI systems. This could result in examples, explanations, names, cultural references, or learning scenarios that unintentionally favor certain regions, languages, or cultural perspectives while underrepresenting others.

There is also a risk that AI-based educational tools may not equally support students from diverse linguistic backgrounds or those with different learning styles. If AI-generated explanations are mainly shaped around dominant cultures or educational systems, some students may find the content less relatable or effective.

However, the audit did not identify evidence of actual representational harm occurring on Khan Academy. These concerns are mainly **structural risks associated with deploying large-scale AI learning systems**, especially for a global student audience. Addressing  e issues requires continuous monitoring of AI outputs, diverse training data, accessibility improvements, and inclusive design practices.

Prompt 5:

How does Khan Academy handle the AI-tutor feature Khanmigo — specifically, who has visibility into student chat logs, and what are the ethical implications for student privacy and psychological safety? Also evaluate the platform's data-security measures: what encryption standards, audit certifications, or third-party assessments does Khan Academy report, and how credible are these claims?

Khan Academy's AI tutor Khanmigo introduces a new category of educational data because students are no longer only completing exercises or watching lessons — they are having conversations with an AI system. According to Khan Academy's policies, Khanmigo chat logs and flagged interactions may be visible to parents, guardians, teachers, or school administrators when a student account is connected through a family or school environment.

This visibility is designed as a **student safety measure**. Adult access allows monitoring of inappropriate, harmful, or concerning interactions and helps schools supervise how AI is being used in learning environments. For younger students especially, this oversight can help identify misuse or safety risks.

However, this creates an important ethical tension around **student privacy and psychological safety**. Students may use an AI tutor to ask questions they feel embarrassed to ask in class, such as repeatedly asking for help on a topic they do not understand. If they know that teachers or parents can later review these conversations, they may avoid asking honest questions or exploring sensitive topics. This reduces the sense of privacy that can make AI tutoring valuable.

There is also a **data power imbalance** between students and institutions. Schools and teachers can access detailed information about learning behavior, progress, and AI interactions, while students may have limited control over how this information is interpreted. Continuous visibility could unintentionally create a feeling of surveillance rather than support.

Regarding **data security**, Khan Academy reports several protective measures:

- **Encryption:** Khan Academy states that it uses modern encryption practices, including **TLS 1.3 encryption for data in transit** and **AES-256 encryption for stored data**. It also applies additional encryption protections for personal data stored through its cloud infrastructure.
- **Security audits:** Khan Academy reports undergoing regular **SOC 2 Type II audits**, which are independent assessments evaluating security controls and data-handling processes.
- **Third-party infrastructure:** The platform uses cloud infrastructure providers such as Google Cloud, meaning some data processing depends on external systems. However, Khan Academy applies its own security controls on top of this infrastructure.

These claims are generally credible because SOC 2 Type II certification involves external auditing rather than only internal promises, and TLS/AES encryption standards are widely accepted security practices. However, encryption and audits do not remove all ethical concerns. They mainly address whether data is protected from unauthorized access, not broader issues such as whether enough data is collected, who should have visibility into AI conversations, or how student autonomy should be balanced with safety.

Verification Table

The table below cross-checks each major AI-generated observation against the actual Khan Academy Privacy Policy, Security Practices page, Children's Privacy Notice, and independent third-party reporting.

AI Observation	Correct / Incorrect	Evidence from Privacy Policy	Missing Ethical Issues	AI Limitations / Hallucinations
AI stated Khan Academy collects only minimal data (name, age, email) and nothing else.	Partially Incorrect	Privacy Policy lists contact/profile info, account & authentication info, location info, browser/device info, donor/applicant info, and usage & learning-activity data (Khan Academy Help Center, 'Security Practices').	AI did not mention device/browser fingerprinting or location data collection.	AI generalized from a typical ed-tech policy template rather than the specific, more detailed clause list.
AI claimed Khan Academy sells or shares data with advertisers for targeted ads.	Incorrect (Hallucination)	Khan Academy explicitly states it does not sell personal information and does not display third-party advertising; it is a Student Privacy Pledge signatory.	N/A	Classic hallucination — AI assumed a common 'freemium' monetization pattern that does not apply to this nonprofit platform.
AI identified that consent is largely bundled ('take it or leave it') via a single checkbox at signup, with no granular opt-in for analytics vs. core functionality.	Mostly Correct	Cookie Notice distinguishes 'strictly necessary' cookies (required) from analytics/marketing cookies (optional), but the general Terms/Privacy Policy acceptance at signup is a single all-or-nothing checkbox.	AI did not note that a cookie-preference banner does allow some granularity, so consent is not fully bundled at the cookie level, only at the account-creation level.	AI oversimplified; nuance between account-level and cookie-level consent was missing.
AI flagged that children under 13 may be using the platform, raising COPPA-related concerns and requiring parental/school consent.	Correct	Khan Academy Kids Privacy Policy and Children's Privacy Notice describe parent/teacher/school-mediated accounts for children, restricted data collection, and no behavioral advertising to children.	AI did not fully explain the school-official-consent model (schools can consent on behalf of parents under FERPA/COPPA), which is central to how Khan Academy actually operationalizes child consent.	AI's answer was accurate at a high level but lacked the specific legal mechanism used (school consent).

AI Observation	Correct / Incorrect	Evidence from Privacy Policy	Missing Ethical Issues	AI Limitations / Hallucinations
AI stated Khan Academy uses 'dark patterns' such as fake urgency and hidden unsubscribe buttons similar to commercial apps.	Incorrect (Hallucination)	No evidence of countdown timers, fake scarcity, or hidden unsubscribe mechanisms was found on the platform or in the policy; email preference and unsubscribe controls are described as accessible.	N/A	AI over-applied dark-pattern taxonomy from commercial/e-commerce apps without verifying against the actual platform.
AI noted that Khanmigo (AI tutor) chat logs are visible to teachers/parents, which could reduce students' sense of psychological safety when seeking help.	Correct	Khan Academy Help Center confirms parents, guardians, teachers, or administrators have visibility into Khanmigo chat logs and flagged interactions, and OpenAI is not permitted to train on student data.	AI did not discuss the moderation/flagging system for self-harm or harmful content as a countervailing safety benefit.	Balanced but one-sided; did not weigh the safety rationale for adult visibility.
AI claimed digital-inequality concerns are irrelevant since Khan Academy is free for everyone.	Incorrect / Oversimplified	Free access does not eliminate the 'digital divide' — device ownership, broadband access, and data costs (relevant in India and other developing regions) still gate access; usage analytics also require sustained connectivity, disadvantaging low-bandwidth users.	AI did not consider infrastructural inequality (device/internet access) or algorithmic personalization bias toward learners with more usage data.	AI equated 'free of cost' with 'equitable access,' a common but flawed assumption.
AI stated Khan Academy encrypts all user data and undergoes independent security audits.	Correct	Khan Academy states it conducts SOC 2 Type 2 audits annually, uses TLS 1.3/AES-256 encryption, and adds an additional encryption layer for personal data on Google Cloud Platform.	None significant.	None — this was a well-grounded, verifiable claim.

Section 5: Ethical Reflection

This audit clarified that responsible data practice is not simply about legal compliance but about designing systems where users can meaningfully understand and control what happens to their information. Khan Academy's nonprofit structure, refusal to sell data, and Student Privacy Pledge commitment show that ethical intent can be embedded in institutional design, not just policy language. Yet even a well-intentioned platform relies on bundled, click-wrap consent that few users read in full, revealing a structural gap between legal consent and genuinely informed consent. True informed consent online is difficult to achieve: policies are long, technical, and presented at the moment of highest motivation to proceed (signup), which discourages careful reading. For minors, the gap widens further, since consent is delegated to parents or schools who may themselves lack full visibility into how AI features like Khanmigo process conversational data. Organizations bear responsibility not only for securing data technically (as Khan Academy does through encryption and SOC 2 audits) but also for designing consent flows that are granular, layered, and revisitable, and for proactively addressing structural harms such as the digital divide and algorithmic bias in AI-generated content, which are easy to overlook when 'the platform is free.' Finally, this exercise showed that Generative AI tools are useful first-pass auditors — fast at summarizing stated policy text — but unreliable as sole authorities. They occasionally hallucinate practices (such as ad-targeting or dark patterns) that do not exist, and they systematically underweight structural and equity-related harms unless explicitly prompted. This reinforces that Generative AI should be used to accelerate, not replace, human critical verification in any ethical audit, especially where real user trust and vulnerable populations (children, low-resource learners) are involved.

References

- Khan Academy Privacy Policy — <https://www.khanacademy.org/about/privacy-policy>
- Khan Academy Kids Privacy Policy — <https://www.khanacademy.org/kids/privacy-policy>
- Khan Academy Help Center, 'What are Khan Academy's Security Practices?' — <https://support.khanacademy.org/hc/en-us/articles/4406593496077>
- Khan Academy Help Center, 'What are Khan Academy's Privacy Principles?' — <https://support.khanacademy.org/hc/en-us/articles/202487520>
- Khan Academy Privacy Principles and Policy (Educator Resource) — <https://www.khanacademy.org/khan-for-educators/resources/teacher-essentials/safety-privacy-and-additional-resources/a/khan-academy-privacy-principles-and-policy>
- Public Services Alliance, 'Student Data Privacy and Security: Khan Academy's Khanmigo' (Oct 2025) — <https://publicservicesalliance.org/2025/10/22/student-data-privacy-and-security-khan-academy-khanmigo/>
- Student Privacy Pledge — <https://studentprivacypledge.org/faqs/>
- Generative AI Tool used for prompts: [ChatGPT / Google Gemini / Microsoft Copilot — specify model/version and access date]